MIPS Simulator: Data Cache
Project 5 – CS 3339 – Fall 2022

Due Date per Canvas.   Friday before 11:59pm, 24-hour late period -10%. After that no submissions accepted.

PROBLEM STATEMENT

In this project, you will further enhance your simulator to model pipeline stalls due to memory latency, and you will simulate a data cache to study the performance impact of caching. You should begin with a copy of your Project 4 submission. Additionally, I have provided a class skeleton for a CacheStats class intended to be instantiated inside your existing CPU class. You may add any functions, function parameters, etc. to it that you want.

The simulated data cache should store 1 KiB (1024 Bytes) of data in block sizes of 8 words (32 bytes). It should be 4-way set associative, with a round-robin replacement policy and a write policy of write-back write-allocate. All blocks in the cache are initially invalid. For simplicity, assume the cache has no write buffer: a store must be completely finished before the processor can proceed. Since this is a data cache, only loads and stores access it; instruction fetches should still be assumed to hit in a perfect I-cache with immediate access (i.e., there is never a stall for an instruction fetch).

Note that since you only have to model the hit/miss/timing behavior of the cache, you do not have to actually store any *data* in your cache model. It is sufficient to simulate the valid, tag, and dirty bits as well as the round-robin replacement policy.

Your cache model will calculate and report the following statistics:
* The total number of **accesses**, plus the number that were **loads** vs. **stores**
* The total number of **misses**, plus the number caused by **loads** vs. **stores**
* The number of **writebacks**
* The **hit ratio**

Every time an access is made to the cache model, the cache model should return the number of cycles that the processor must stall in order for that access to complete. It takes 0 cycles to do a lookup or to hit in the cache (i.e., data that is hit will be returned or written immediately). A read access to the next level of the memory hierarchy (e.g., main memory) has a latency of 30 cycles, and a write access has a latency of 10 cycles. Note that an access resulting in the replacement of a dirty line requires both a main memory write (to write back the dirty block) and a read (to fetch the new block) consecutively. Because the cache has no write buffer, all stores must stall until the write is complete.

Before computing the cache statistics, be sure to drain all the dirty data from the cache; i.e., write it back. Count these as writebacks, but do not count any stalls/latency resulting from them.

Inside your Stats class, add a new function similar to the bubble() and flush() functions. This stall() function should stall the entire pipeline for a specified number of cycles. The Stats class should track the total number of stall cycles that occur during program execution.

Your simulator will report the following statistics at the end of the program:
- The exact **number of clock cycles** it would take to execute the program on this CPU
- The **CPI** (cycle count / instruction count)
- *The **number of bubble cycles** injected due to data dependencies (unchanged from Proj. 4)*
- *The **number of flush cycles** in the shadows of jumps and taken branches (also unchanged)*
- The **number of stall cycles** due to cache/memory latency
- The **data cache statistics** reported by the cache model

I have provided the following new files:
- A `CacheStats.h` class specification file, to which you'll need to add member variables
- A `CacheStats.cpp` class implementation file, which you should enhance with code to model the described cache and count accesses, misses, and writebacks
- A new Makefile

In addition to enhancing the `CacheStats.h/.cpp` skeleton, you will need to modify your existing `Stats.h/Stats.cpp` in order to implement memory stalls. You will also need to modify `CPU.h` to instantiate a CacheStats object, and `CPU.cpp` to call both CacheStats and Stats class functions appropriately to model cache behavior and resulting pipeline stalls. You'll also need to change `CPU::printFinalStats()` to match my expected output format (see below).

ASSIGNMENT SPECIFICS

Begin by copying all of your Project 4 files into a new Project 5 directory. Then add the additional files from Canvas to your project5 directory. You can compile and run the simulator program identically to previous projects, and test it using the same `*.mips` inputs. (Only `sssp.mips` yields interesting cache behavior; I'll only test your code with this one).

If you examine `CacheStats.h`, you'll notice that I've already defined constants for you for all of the cache configuration options you'll need (e.g., number of sets, number of ways, block size, read miss latency, etc.). You should not need to change any of these defines. They are defined to be modifiable from the compilation command line, but for this project, I will not change any of them. You can even get away with not using these defined constants if you prefer.

The following is the expected result for sssp.mips. Your output must match this format verbatim:

```
CS 3339 MIPS Simulator
Cache Config: 1024 B (32 bytes/block, 8 sets, 4 ways)
  Latencies: Lookup = 0 cycles, Read = 30 cycles, Write = 10 cycles
Running: sssp.mips

 7 1


Program finished at pc = 0x400440  (449513 instructions executed)

Cycles: 1396800
CPI: 3.11

Bubbles: 481710
Flushes: 51990
Stalls:  413580

Accesses: 197484
  Loads: 146709
  Stores: 50775
Misses: 12044
  Load misses: 8559
  Store misses: 3485
Writebacks: 5229
Hit Ratio: 93.9%
```

The CacheStats skeleton already includes code to disable the cache (i.e., all loads result in a read access to the next level of the memory hierarchy, and all stores result in a write access). To explore the performance impact of adding a cache, you can compile your simulator with the cache disabled:

```
$ make clean; make CACHE_EN=0
```

Re-run the simulator on sssp.mips. What happens to the CPI? How big is the difference?

Additional Requirements:
- **Your code must compile with the given `Makefile` and run on zeus.cs.txstate.edu**
- Your code must be well-commented, sufficient to prove you understand its operation
- Make sure your code doesn't produce unwanted output such as debugging messages. (You can accomplish this by using the `D(x)` macro defined in `Debug.h`)
- Make sure your code's runtime is not excessive
- Make sure your code is correctly indented and uses a consistent coding style
- Clean up your code before submitting: i.e., make sure there are no unused variables, unreachable code, etc.

SUBMISSION INSTRUCTIONS

Submit all of the code necessary to compile your simulator (all of the `.cpp` and `.h` files as well as the `Makefile`) as a compressed tarball. You can do this using the following Linux command:

```
$ tar czvf yourNetID_project5.tgz *.cpp *.h Makefile
```

Do not submit the executables (`*.mips` files). Any special instructions or comments to the grader, including notes about features you know do not work, should be included in a separate text file (not inside the tarball) named `yourNetID_README.txt`.

Use the submit_test script that was provided for assignment 3 to confirm the contents of your tar file, ability to compile and execute on zeus, and check your output. You will need to copy over the input sssp.mips and use the new version of sssp.out provided with Project 5. The remaining instructions are in the script itself. ***Submissions that do not untar correctly and compile on zeus with this script will receive zero points, so you are highly encouraged to check your submission.***

As in the prior projects and assignments all files are to be submitted using Canvas. Note that files are only submitted if Canvas indicates a successful submission. Pictures of timestamped files on your local machine or zeus will not be accepted as proof of work completed.

You may submit your file(s) as many times as you'd like before the deadline. Only the last submission will be graded. ***Canvas will not allow submission after the deadline***, so I strongly recommend that you don't come down to the final seconds of the assignment window. After Canvas closes, no submissions will be accepted.

ACADEMIC HONESTY (excerpt from course syllabus)

You are expected to adhere to the Texas State University Honor Code which is described here
http://www.txstate.edu/honorcodecouncil/Academic-Integrity.html

All assignments and exams must be done individually, unless otherwise noted in the assignment handout. Turning in an exam or assignment that is not entirely your own work is cheating and will not be tolerated.

Group discussion about course content is NOT cheating and is in fact strongly encouraged! You are welcome to study together and to discuss information and concepts covered in class, as well as to offer (or receive) help with debugging or understanding course concepts to (or from) other students. However, this cooperation should never involve students possessing a copy of work done by another student, including solutions from previous semesters, other course sections, the Internet, or any other sources.

Turning in an assignment any part of which is copied from the Internet, another student's work, or any other non-approved source will result in a 0 grade on the assignment and will be reported to the Texas State Honor Code Council. Should one student copy from another, both the student who copied work and the student who gave material to be copied will receive 0s and be reported to the Honor Code Council. You should never grant anyone access to your files or email your programs to anyone (other than the instructor)!